

The reproducibility crises

02457 Machine Learning Operations

Nicki Skafte Detlefsen,

Postdoc

DTU Compute

What is it?



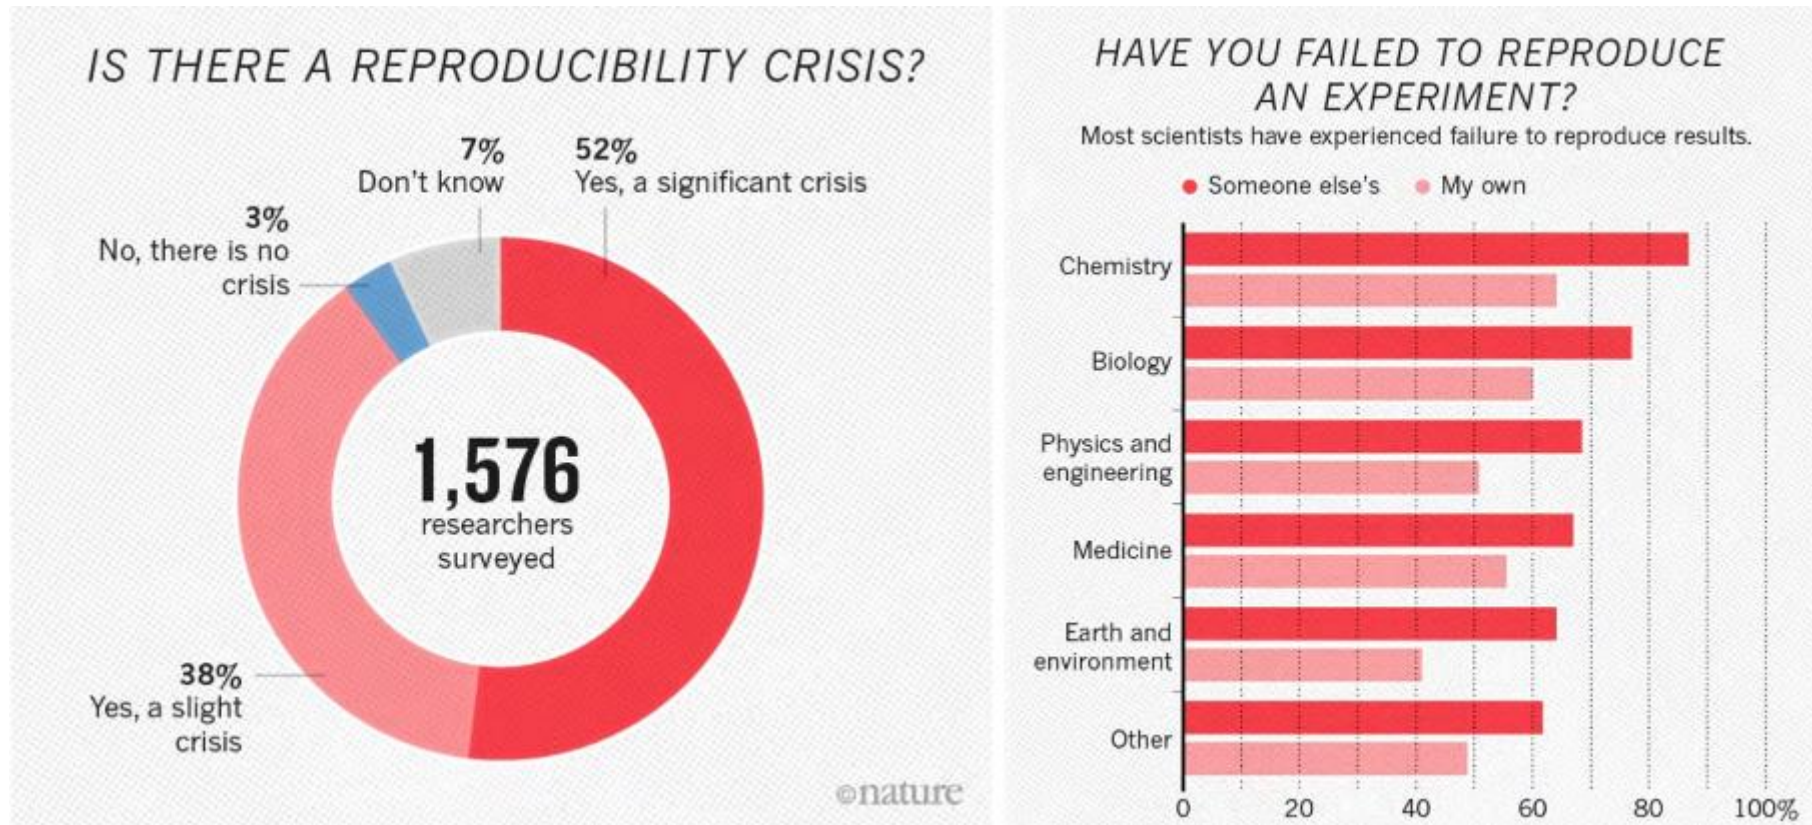
- Being able to reproduce other peoples experimental results is an essential part of the scientific method
- Well known problem throughout most fields (physics, chemistry, biology and computer science)
- With the rise of deep learning, the problem has only been made worse due to competition



This is where it breaks

How bad is it?

Wow its bad...




<https://www.nature.com/news/1-500-scientists-lift-the-lid-on-reproducibility-1.19970>

What can we do about it

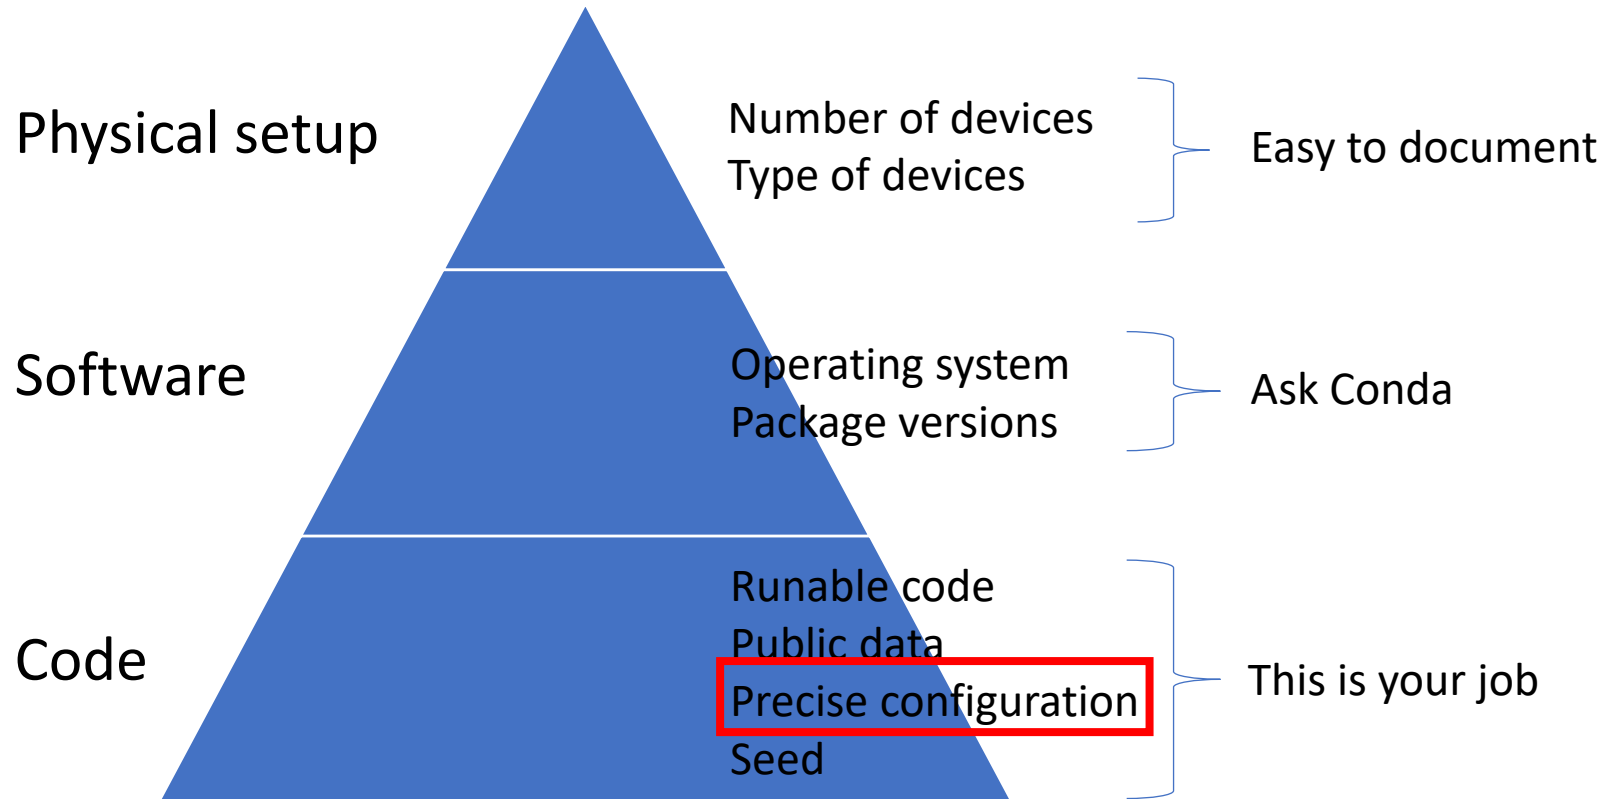


- Make sure to document everything about our experiments



Document this step
as thoroughly as
possible

Closer look



Configurations



There is a lot of subjective choices that we do when running experiments, most notable the hyperparameters.

Parameters in script	Argparser	Config files
<pre>class hparams: lr = 0.1 batch_size = 16 num_layers = 5</pre>	<pre>python my_script.py \ --lr 0.1 \ --batch_size 16 \ --num_layers 5</pre>	<pre>experiment1.yaml lr: 0.001 batch_size: 16 num_layers: 5 python my_script.py \ config=experiment1.yaml</pre>
Not easy to configure Experiments may be lost if not careful	Easy to configure Falls on user to save the configuration	Easy to configure Parameters are systematically saved with the experiment

Hydra



A framework for elegantly configuring complex applications.

<https://github.com/facebookresearch/hydra>



Other frameworks



- <https://github.com/IDSIA/sacred>
- <https://mlflow.org/>
- <https://wandb.ai/site>